

Sat 5 – Sun 6 September | Doors 9:00am

# Rebuild private markets

Ylookup × Encode AI Hackathon — let's bring tech to a dinosaur industry.

Encode Hub, Shoreditch · Come solo or as a team of 4 · \$2,000 prize pool



**Join the Discord**

Scan to get announcements, mentors and support

FIRST THING, RIGHT NOW

# Join the Discord



Announcements, mentor requests, and everything you need all weekend live there. Scan it before you do anything else.

## SCHEDULE

# Saturday

09:00 – 10:00

Doors open · check-in

10:00 – 11:00

**Welcome & intro** — these slides, the domain briefing and both track briefs

11:00 – 12:00

Team forming — mingle, find your team, pick your track

12:00

**Hacking begins.**

13:00

**Guest talk - Simon from Thinking Machines Lab** · lunch served

19:00

Dinner

Late

Keep building. Venue access details at the desk

## SCHEDULE

# Sunday

Submissions close 12:00 — hard stop

09:00

Doors open · breakfast

11:00

**Final push** · record and upload your demo video

12:00

**Submissions close.** Tally form shuts, repos are frozen

12:30

Lunch

13:30

Judging round · live demos to the judging panel

16:00

Winners announced

16:30

Close

THE INDUSTRY YOU ARE BUILDING FOR


# \$30T

**Private markets manage thirty trillion dollars in assets. The machinery that keeps the books is still run almost entirely by hand.**

A high-level look at where it breaks is next. The specifics are yours to find.

## BACKGROUND

# Money goes down. Numbers have to come back up.



Every arrow on this diagram is money moving. Each one has to be recorded, valued and reported — otherwise nobody knows what anything is worth.

## THE INVESTOR ASKS

"What is my stake worth today?"

"How much have I actually made?"

The manager owes them that answer, but wants to spend the day investing — not reconciling. So they outsource it.

**Fund admins are the back office of funds — they do the bookkeeping for all of this: valuation, accounting, financial statements. Essentially an accountant for funds.**

Valuations

Bookkeeping

Accounting

Financial statements

THE HIGH-LEVEL PICTURE

# Old technology, and more people — instead of better process.

## It runs on decades-old technology

The systems were never rebuilt around software, and they were never made to talk to each other.

## So growth is answered with headcount

More clients means more people doing the same manual work — rather than better process or genuine innovation.

That is as far as we take you. What actually breaks because of it is in the interviews — working that out is the point of the weekend.

## THE TRACKS

# Pick one. Both win the same prizes.

### PRODUCT TRACK

**Build a product around a real problem in private capital.**

We hand you anonymised fund manager interviews and a fund's worth of data. You find the problem worth solving, then build against it.

Engineers

Designers

Non-technical welcome

### RESEARCH TRACK

**Push an AI model to solve complex fund workflows.**

Harnesses, fine-tuning, prompting, evals — method open. Full brief on the next slides.

AI researchers

ML engineers

TRACK ONE OF TWO

# Product track

---

The problem-finding side: talk to the industry, pick a real pain point, build for it.

Anonymised interviews · live personas · your product

# How the track works

## 01 · READ THE INTERVIEWS

Anonymised transcripts of **real user interviews with fund managers** — us asking them where fund administration hurts.

Their pain points, in their own words.  
That is your raw material.

## 02 · INTERVIEW US

From **11:30 to 13:00**, three of us are open for questions — one playing each side of the problem.

Ask anything about the industry.  
Details on the next slide.

## 03 · PICK ONE AND BUILD

Method open — any stack, any approach. We brief a problem space, not a spec.

Tell us which problem you picked, and why.

**We do not hand you the answer. What a hundred outsiders see in this research that we have not is the whole point.**

## PRODUCT TRACK

# Talk to the industry, 11:30 to 13:00

Three of us from the Ylookup team are open for questions — each speaking as one side of the problem, so you can see it from every angle.

### PERSONA 01

## The fund manager

The customer. Ask what they need from their admin, what they never get, and what they do not trust.

### PERSONA 02

## The fund admin

The service provider. Ask how the work is organised, priced and handed down the chain.

### PERSONA 03

## The fund accountant

The person doing it. Ask what the day actually looks like, tool by tool, spreadsheet by spreadsheet.

The transcripts tell you what hurts. These conversations tell you why it is still like that.

# Four criteria. 25% each.

Two on judgement, two on craft. None carries more than another, and none is optional.

## 25% Problem identification

Did you find a real problem in the interviews — and can you show us where it came from?

## 25% Product

Does what you built actually solve the problem you chose to tackle?

## 25% UI

Clean and considered. A non-technical fund manager is the user. No AI slop.

## 25% Code review

Structure and style a working team could live with. Reviewed on submission.

# How to submit

## 01

### Tally form

One submission per team, through the Tally form we share at team-forming. Link also pinned in the event channel.

## 02

### Public GitHub repo

Public at submission — the code review is scored from it. At minimum it must **run from a script**: README plus one command we can execute.

## 03

### 3–5 min demo video

Open with the problem you picked and the interview it came from. Then show the product working.

Hosting is optional, but it earns extra points.

Deadline Sunday 12:00. Submit early, resubmit if you need to.

TRACK TWO OF TWO

# Research track

---

The raw accuracy side: make a small model do real spreadsheet work.

One benchmark · one model · your architecture

# Workbook in. Workbook out. Cells must match.

You get a spreadsheet and a plain-English instruction written by someone on an Excel forum. You return the same workbook with the answer filled in. **SpreadsheetBench Verified — 400 human-validated tasks.**

## ONE TASK

Instruction

"How can I create a VBA code that will count..."

Given

1\_51-12\_init.xlsx · data at A1:M7

Graded

Sheet1!B6 against the golden workbook

## WHAT GETS COMPARED

275 cell-level · 125 sheet-level

Your output

Golden workbook

|   | A     | B  |   |   | A     | B  |
|---|-------|----|---|---|-------|----|
| 4 | Mar   | 18 | ↔ | 4 | Mar   | 18 |
| 5 | Apr   | 31 |   | 5 | Apr   | 31 |
| 6 | Total | 49 |   | 6 | Total | 49 |
|   |       |    |   |   |       |    |

Only the graded cells are diffed — **B6 here**. Everything else in the workbook is ignored, formatting included.

# The caveat: a very small model

## Qwen3.8-27B

Same model for every team. No frontier models.

It scores **59.0% pass** one-shot on the 400. That is the number you start from — everything you win this weekend comes from what you build around it, not from a bigger model.

### FINE-TUNING IS ON THE TABLE

## Tinker, from Thinking Machines Lab

We are giving you access. Train a checkpoint and run the same harness against it — accounts on the next slide.

### METHOD IS WIDE OPEN

## No prescribed approach

Choose whatever agentic architecture you want: tool loops, code execution, self-checking, retries, retrieval, or plain prompting.

### AND YOU START WITH

A working repo — `uv sync`, then the 400 tasks downloaded

A one-shot baseline you can beat immediately

The official scorer the judges run

# Every graded cell, or the task fails

## THE RANKING

### **pass\_rate**

Share of the tasks where **every** graded cell matches the golden workbook, after recalculation. One wrong cell fails the task.

## THE TIE-BREAK

### **cell\_accuracy**

Graded cells matched across all tasks — plus the pass rate on each instruction type, so we can see where you are strong.

ALL 400  
COUNT

Judges score the whole set. A missing prediction is a failed task — never skip ids to protect a number.

# The score is the main thing. It is not the only thing.

## WHAT WE WEIGH MOST

### Pass rate, then cell accuracy

Optimise for this. It is the number that ranks the track, and nothing below outweighs it.

## WE ALSO TEST ON

### A private dataset of real fund data

We also score the graded cells on a held-back set of real fund spreadsheets you never see. Something tuned only to SpreadsheetBench will show up here.

## AND WE READ

### Your approach

The write-up is read by the judging panel and scored on how original and well-reasoned the method is — what you tried, and why.

Judged on approach and results together — so write down what you did. An interesting method loses points if nobody can tell what it was.

# Two accounts, in this order

## 01

### A Google Cloud account

One per team member. Use an account you control — a locked-down work account often will not work.

Give us that address at the research desk — it is what the invite below is sent to.

## 02

### Thinking Machines Lab

**We invite you to the platform once you have formed your team and picked the research track.** You sign in with Google — that is what unlocks fine-tuning with Tinker.

And **Simon Guo**, a researcher at Thinking Machines Lab, gives a guest talk later — he will show you what the platform can do.

Invites are sent manually, so get your address to us at team-forming — not at 3pm when you want to launch a training run.

# How to submit

You ship a container, not a notebook. We run it against a **holdout set you have never seen**, mounted read-only — so it has to work unattended, first time.

## KEEP THE CONTRACT

Start from the reference `Dockerfile`, swap in your pipeline

Holdout read-only at `/data`, results to `/out`

Four output files: `predictions.jsonl`, `outputs/`, `traces/`, `run.log`

Model-written code runs inside the container

## AND HAND IN

A **public repo** — same code-review criterion as the product track

Your own scored run on the 400, so we can see what you expected

**Required:** a method write-up — what you tried, what moved the number

Full details in `SUBMISSION.md`

**Deadline: Sunday 12:00, same as the product track. Build the image early and test the run — a container that does not start scores zero.**

## RULES

# Teams & fair play

---

### Solo, or up to 4

Team forming runs 11:00–12:00 Saturday. We help match people who arrive alone.

---

### Brand new ideas

We encourage something built fresh this weekend. Boilerplate and open-source libraries are fine — an existing project is not.

---

### Be decent

Encode's code of conduct applies. Anything uncomfortable, tell an organiser.

---

### One track per team

Declare it when you form. Switching after 13:00 Saturday is not practical.

---

### AI tooling encouraged

Use whatever helps. You still own the architecture — code review assumes you can defend it.

## PRIZES & OPPORTUNITIES

# What you walk away with

---

### **\$2,000 prize pool**

Apple Watches, AirPods and more cool hardware and accessories.

---

### **Pitch EQT Ventures**

Winners pitch the \$12B fund backing us.

---

Both tracks are judged for the same prizes.

# Everything you need, free for the weekend

## Model credits

Build credits for every team, issued at team-forming. Ask at the desk if you burn through them.

## API keys

One key per team. Do not commit it to your public repo — judges will notice.

## Data pack & interviews

Everything is on the Discord server already — released from the start, no waiting.

## Venue

Encode Hub, Shoreditch. Wifi, power and food on site — details on the door.

# Ask us anything. Especially about the industry.

## Product track & the data


Ylookup engineers, on the floor all weekend. Grab any of us whenever you need to.

## How funds actually work

A real fund manager is on site Saturday afternoon. Use them — nobody else gets this.

## Logistics & the venue

Encode organisers at the front desk, plus the event channel for anything async.



Hacking begins Sat 12:00 · Submissions close Sun 12:00 · Winners 16:00



Join the Discord